

Inference latency vs T on NVIDIA A100 — grouped bars (log-y)

