

Inference latency vs T on NVIDIA A100

